

```
In [91]: import os
%matplotlib inline
# Prevent CUDA from using GPU as it does not work well on my pc
os.environ["CUDA_VISIBLE_DEVICES"] = "-1"

# Set Constants of the model
BATCH_SIZE = 64
SHUFFLE_BUFFER_SIZE = 100
```

```
In [2]: # Helper functions
import numpy as np

# Breaks down a list of integer values into a one-hot like format
def one_hot_training(np_array):
    transformed_list = []
    for arr in np_array:
        new_arr = np.zeros(10)
        new_arr[int(arr)] = 1
        transformed_list.append(new_arr)
    return np.array(transformed_list)

# This translates the highest value from the one-hot encoding into the correct sign name
def one_hot_translator(np_array):
    labels_names = ['Stop', 'Yield', 'Red Light', 'Green Light', 'Roundabout', 'Right Turn O',
                    'Do Not Enter', 'Crosswalk', 'Handicap Parking', 'No Parking']
    return labels_names[np.argmax(np_array)]

# This translates an entire array of one-hot encoded sign predictions
def translate_all(np_array):
    translated_values = []
    for i in np_array:
        translated_values.append(one_hot_translator(i))
    return np.array(translated_values)
```

```
In [3]: # First import the data
import tensorflow as tf
data_train = np.load('data_train.npy').transpose()
labels_train = np.load('labels_train.npy')
data_train = np.array([i.reshape(300,300,3) for i in data_train])
data_train = np.array(tf.cast(tf.image.resize(data_train,(150,150)), np.uint8))
```

```
In [4]: # Process the data so that it is in the expected form for the InceptionV3 model
import tensorflow as tf
processed = tf.keras.applications.inception_v3.preprocess_input(data_train, data_format

# Break down data into training and test sets
from sklearn.model_selection import train_test_split
x_train, x_test, t_train, t_test = train_test_split(processed, one_hot_training(labels_
```

```
In [5]: # Augment data to reduce overfitting

from tensorflow.keras.preprocessing.image import ImageDataGenerator
train_datagen = ImageDataGenerator(horizontal_flip=True,
```

```

        vertical_flip=True,
        rotation_range=90,
        brightness_range=(.75, 1))

train_generator = train_datagen.flow(
    x_train,
    y = t_train,
    batch_size=BATCH_SIZE)

```

In [6]:

```

# Import the InceptionV3 Model

from tensorflow.keras.applications.inception_v3 import InceptionV3
inception = InceptionV3(input_shape=(150,150,3),
                        include_top=False,
                        weights='imagenet')

# Set layers to false to prevent overwriting the existing model
for layer in inception.layers:
    layer.trainable = False

# Create output layers that will be trained
from tensorflow.keras.optimizers import SGD
x = tf.keras.layers.Flatten()(inception.output)
x = tf.keras.layers.Dense(1024, activation="relu")(x)
x = tf.keras.layers.Dropout(0.15)(x)
x = tf.keras.layers.Dense(10, activation='softmax')(x)

# Create Optimizer
Adam = tf.keras.optimizers.Adam(learning_rate=0.001, beta_1=0.9, beta_2=0.999, epsilon=
Nadam = tf.keras.optimizers.Nadam(learning_rate=0.001, beta_1=0.9, beta_2=0.999, epsilo
SGD = SGD(learning_rate=0.01, nesterov=True)
optimizer = SGD

# Finalize and compile the model
model = tf.keras.Model(inception.input, outputs = x)
model.compile(optimizer = optimizer,
              loss = 'categorical_crossentropy',
              metrics = ['categorical_accuracy', 'acc', 'mean_squared_error'])

```

In [7]:

```

# Fit the model to the dataset
es = tf.keras.callbacks.EarlyStopping(monitor='acc', mode='max', verbose=1, patience=10)
callbacks = tf.keras.callbacks.Callback()
history = model.fit(train_generator, epochs=140, batch_size=BATCH_SIZE, callbacks=[es])
model.save("140_epoch.h5")

```

```

Epoch 1/140
78/78 [=====] - 40s 475ms/step - loss: 2.1024 - categorical_acc
uracy: 0.6959 - acc: 0.6959 - mean_squared_error: 0.0451
Epoch 2/140
78/78 [=====] - 38s 487ms/step - loss: 0.4627 - categorical_acc
uracy: 0.8565 - acc: 0.8565 - mean_squared_error: 0.0206
Epoch 3/140
78/78 [=====] - 37s 472ms/step - loss: 0.3757 - categorical_acc
uracy: 0.8931 - acc: 0.8931 - mean_squared_error: 0.0164
Epoch 4/140
78/78 [=====] - 36s 464ms/step - loss: 0.3250 - categorical_acc
uracy: 0.9036 - acc: 0.9036 - mean_squared_error: 0.0144

```

Epoch 5/140  
78/78 [=====] - 36s 461ms/step - loss: 0.3092 - categorical\_acc  
uracy: 0.9118 - acc: 0.9118 - mean\_squared\_error: 0.0132  
Epoch 6/140  
78/78 [=====] - 36s 462ms/step - loss: 0.2946 - categorical\_acc  
uracy: 0.9151 - acc: 0.9151 - mean\_squared\_error: 0.0127  
Epoch 7/140  
78/78 [=====] - 37s 467ms/step - loss: 0.2733 - categorical\_acc  
uracy: 0.9221 - acc: 0.9221 - mean\_squared\_error: 0.0118  
Epoch 8/140  
78/78 [=====] - 37s 468ms/step - loss: 0.2688 - categorical\_acc  
uracy: 0.9249 - acc: 0.9249 - mean\_squared\_error: 0.0116  
Epoch 9/140  
78/78 [=====] - 37s 466ms/step - loss: 0.2391 - categorical\_acc  
uracy: 0.9326 - acc: 0.9326 - mean\_squared\_error: 0.0104  
Epoch 10/140  
78/78 [=====] - 37s 465ms/step - loss: 0.2465 - categorical\_acc  
uracy: 0.9334 - acc: 0.9334 - mean\_squared\_error: 0.0104  
Epoch 11/140  
78/78 [=====] - 36s 464ms/step - loss: 0.2149 - categorical\_acc  
uracy: 0.9383 - acc: 0.9383 - mean\_squared\_error: 0.0095  
Epoch 12/140  
78/78 [=====] - 37s 469ms/step - loss: 0.2099 - categorical\_acc  
uracy: 0.9411 - acc: 0.9411 - mean\_squared\_error: 0.0092  
Epoch 13/140  
78/78 [=====] - 37s 467ms/step - loss: 0.2060 - categorical\_acc  
uracy: 0.9407 - acc: 0.9407 - mean\_squared\_error: 0.0090  
Epoch 14/140  
78/78 [=====] - 37s 465ms/step - loss: 0.2102 - categorical\_acc  
uracy: 0.9403 - acc: 0.9403 - mean\_squared\_error: 0.0092  
Epoch 15/140  
78/78 [=====] - 37s 465ms/step - loss: 0.1916 - categorical\_acc  
uracy: 0.9459 - acc: 0.9459 - mean\_squared\_error: 0.0085  
Epoch 16/140  
78/78 [=====] - 37s 477ms/step - loss: 0.1919 - categorical\_acc  
uracy: 0.9465 - acc: 0.9465 - mean\_squared\_error: 0.0084  
Epoch 17/140  
78/78 [=====] - 37s 467ms/step - loss: 0.1870 - categorical\_acc  
uracy: 0.9473 - acc: 0.9473 - mean\_squared\_error: 0.0082  
Epoch 18/140  
78/78 [=====] - 37s 467ms/step - loss: 0.1913 - categorical\_acc  
uracy: 0.9447 - acc: 0.9447 - mean\_squared\_error: 0.0085  
Epoch 19/140  
78/78 [=====] - 36s 465ms/step - loss: 0.1691 - categorical\_acc  
uracy: 0.9508 - acc: 0.9508 - mean\_squared\_error: 0.0076  
Epoch 20/140  
78/78 [=====] - 37s 465ms/step - loss: 0.1754 - categorical\_acc  
uracy: 0.9483 - acc: 0.9483 - mean\_squared\_error: 0.0080  
Epoch 21/140  
78/78 [=====] - 37s 465ms/step - loss: 0.1823 - categorical\_acc  
uracy: 0.9467 - acc: 0.9467 - mean\_squared\_error: 0.0082  
Epoch 22/140  
78/78 [=====] - 37s 468ms/step - loss: 0.1691 - categorical\_acc  
uracy: 0.9548 - acc: 0.9548 - mean\_squared\_error: 0.0072  
Epoch 23/140  
78/78 [=====] - 37s 465ms/step - loss: 0.1492 - categorical\_acc  
uracy: 0.9580 - acc: 0.9580 - mean\_squared\_error: 0.0067  
Epoch 24/140  
78/78 [=====] - 37s 466ms/step - loss: 0.1664 - categorical\_acc  
uracy: 0.9518 - acc: 0.9518 - mean\_squared\_error: 0.0074

Epoch 25/140  
78/78 [=====] - 37s 468ms/step - loss: 0.1494 - categorical\_acc  
uracy: 0.9586 - acc: 0.9586 - mean\_squared\_error: 0.0066  
Epoch 26/140  
78/78 [=====] - 37s 469ms/step - loss: 0.1433 - categorical\_acc  
uracy: 0.9560 - acc: 0.9560 - mean\_squared\_error: 0.0067  
Epoch 27/140  
78/78 [=====] - 37s 466ms/step - loss: 0.1447 - categorical\_acc  
uracy: 0.9598 - acc: 0.9598 - mean\_squared\_error: 0.0065  
Epoch 28/140  
78/78 [=====] - 37s 468ms/step - loss: 0.1345 - categorical\_acc  
uracy: 0.9607 - acc: 0.9607 - mean\_squared\_error: 0.0062  
Epoch 29/140  
78/78 [=====] - 37s 468ms/step - loss: 0.1432 - categorical\_acc  
uracy: 0.9615 - acc: 0.9615 - mean\_squared\_error: 0.0064  
Epoch 30/140  
78/78 [=====] - 37s 466ms/step - loss: 0.1304 - categorical\_acc  
uracy: 0.9617 - acc: 0.9617 - mean\_squared\_error: 0.0060  
Epoch 31/140  
78/78 [=====] - 37s 468ms/step - loss: 0.1269 - categorical\_acc  
uracy: 0.9637 - acc: 0.9637 - mean\_squared\_error: 0.0058  
Epoch 32/140  
78/78 [=====] - 37s 467ms/step - loss: 0.1252 - categorical\_acc  
uracy: 0.9629 - acc: 0.9629 - mean\_squared\_error: 0.0058  
Epoch 33/140  
78/78 [=====] - 37s 469ms/step - loss: 0.1243 - categorical\_acc  
uracy: 0.9637 - acc: 0.9637 - mean\_squared\_error: 0.0059  
Epoch 34/140  
78/78 [=====] - 37s 469ms/step - loss: 0.1307 - categorical\_acc  
uracy: 0.9592 - acc: 0.9592 - mean\_squared\_error: 0.0060  
Epoch 35/140  
78/78 [=====] - 37s 466ms/step - loss: 0.1208 - categorical\_acc  
uracy: 0.9657 - acc: 0.9657 - mean\_squared\_error: 0.0055  
Epoch 36/140  
78/78 [=====] - 37s 469ms/step - loss: 0.1236 - categorical\_acc  
uracy: 0.9647 - acc: 0.9647 - mean\_squared\_error: 0.0058  
Epoch 37/140  
78/78 [=====] - 37s 467ms/step - loss: 0.1254 - categorical\_acc  
uracy: 0.9631 - acc: 0.9631 - mean\_squared\_error: 0.0058  
Epoch 38/140  
78/78 [=====] - 37s 468ms/step - loss: 0.1191 - categorical\_acc  
uracy: 0.9659 - acc: 0.9659 - mean\_squared\_error: 0.0054  
Epoch 39/140  
78/78 [=====] - 37s 468ms/step - loss: 0.1196 - categorical\_acc  
uracy: 0.9641 - acc: 0.9641 - mean\_squared\_error: 0.0057  
Epoch 40/140  
78/78 [=====] - 37s 470ms/step - loss: 0.1072 - categorical\_acc  
uracy: 0.9675 - acc: 0.9675 - mean\_squared\_error: 0.0050  
Epoch 41/140  
78/78 [=====] - 37s 469ms/step - loss: 0.1116 - categorical\_acc  
uracy: 0.9695 - acc: 0.9695 - mean\_squared\_error: 0.0049  
Epoch 42/140  
78/78 [=====] - 37s 467ms/step - loss: 0.1214 - categorical\_acc  
uracy: 0.9637 - acc: 0.9637 - mean\_squared\_error: 0.0057  
Epoch 43/140  
78/78 [=====] - 37s 469ms/step - loss: 0.1035 - categorical\_acc  
uracy: 0.9643 - acc: 0.9643 - mean\_squared\_error: 0.0050  
Epoch 44/140  
78/78 [=====] - 37s 467ms/step - loss: 0.1029 - categorical\_acc  
uracy: 0.9673 - acc: 0.9673 - mean\_squared\_error: 0.0049


Epoch 45/140  
78/78 [=====] - 37s 469ms/step - loss: 0.1046 - categorical\_acc  
uracy: 0.9679 - acc: 0.9679 - mean\_squared\_error: 0.0049  
Epoch 46/140  
78/78 [=====] - 37s 469ms/step - loss: 0.0994 - categorical\_acc  
uracy: 0.9681 - acc: 0.9681 - mean\_squared\_error: 0.0048  
Epoch 47/140  
78/78 [=====] - 37s 467ms/step - loss: 0.1030 - categorical\_acc  
uracy: 0.9683 - acc: 0.9683 - mean\_squared\_error: 0.0048  
Epoch 48/140  
78/78 [=====] - 37s 466ms/step - loss: 0.1017 - categorical\_acc  
uracy: 0.9707 - acc: 0.9707 - mean\_squared\_error: 0.0048  
Epoch 49/140  
78/78 [=====] - 37s 468ms/step - loss: 0.0911 - categorical\_acc  
uracy: 0.9736 - acc: 0.9736 - mean\_squared\_error: 0.0042  
Epoch 50/140  
78/78 [=====] - 37s 466ms/step - loss: 0.0958 - categorical\_acc  
uracy: 0.9709 - acc: 0.9709 - mean\_squared\_error: 0.0046  
Epoch 51/140  
78/78 [=====] - 37s 466ms/step - loss: 0.0944 - categorical\_acc  
uracy: 0.9740 - acc: 0.9740 - mean\_squared\_error: 0.0044  
Epoch 52/140  
78/78 [=====] - 37s 465ms/step - loss: 0.0918 - categorical\_acc  
uracy: 0.9736 - acc: 0.9736 - mean\_squared\_error: 0.0044  
Epoch 53/140  
78/78 [=====] - 37s 467ms/step - loss: 0.0855 - categorical\_acc  
uracy: 0.9742 - acc: 0.9742 - mean\_squared\_error: 0.0041  
Epoch 54/140  
78/78 [=====] - 37s 467ms/step - loss: 0.0935 - categorical\_acc  
uracy: 0.9734 - acc: 0.9734 - mean\_squared\_error: 0.0044  
Epoch 55/140  
78/78 [=====] - 37s 473ms/step - loss: 0.0882 - categorical\_acc  
uracy: 0.9736 - acc: 0.9736 - mean\_squared\_error: 0.0042  
Epoch 56/140  
78/78 [=====] - 37s 466ms/step - loss: 0.0894 - categorical\_acc  
uracy: 0.9738 - acc: 0.9738 - mean\_squared\_error: 0.0041  
Epoch 57/140  
78/78 [=====] - 37s 469ms/step - loss: 0.0786 - categorical\_acc  
uracy: 0.9770 - acc: 0.9770 - mean\_squared\_error: 0.0038  
Epoch 58/140  
78/78 [=====] - 37s 469ms/step - loss: 0.0829 - categorical\_acc  
uracy: 0.9746 - acc: 0.9746 - mean\_squared\_error: 0.0039  
Epoch 59/140  
78/78 [=====] - 37s 470ms/step - loss: 0.0823 - categorical\_acc  
uracy: 0.9764 - acc: 0.9764 - mean\_squared\_error: 0.0039  
Epoch 60/140  
78/78 [=====] - 37s 472ms/step - loss: 0.0873 - categorical\_acc  
uracy: 0.9762 - acc: 0.9762 - mean\_squared\_error: 0.0040  
Epoch 61/140  
78/78 [=====] - 38s 487ms/step - loss: 0.0788 - categorical\_acc  
uracy: 0.9770 - acc: 0.9770 - mean\_squared\_error: 0.0037  
Epoch 62/140  
78/78 [=====] - 39s 500ms/step - loss: 0.0819 - categorical\_acc  
uracy: 0.9748 - acc: 0.9748 - mean\_squared\_error: 0.0038  
Epoch 63/140  
78/78 [=====] - 39s 499ms/step - loss: 0.0767 - categorical\_acc  
uracy: 0.9784 - acc: 0.9784 - mean\_squared\_error: 0.0034  
Epoch 64/140  
78/78 [=====] - 39s 501ms/step - loss: 0.0722 - categorical\_acc  
uracy: 0.9784 - acc: 0.9784 - mean\_squared\_error: 0.0034

Epoch 65/140  
78/78 [=====] - 38s 488ms/step - loss: 0.0701 - categorical\_acc  
uracy: 0.9792 - acc: 0.9792 - mean\_squared\_error: 0.0034  
Epoch 66/140  
78/78 [=====] - 37s 475ms/step - loss: 0.0741 - categorical\_acc  
uracy: 0.9768 - acc: 0.9768 - mean\_squared\_error: 0.0036  
Epoch 67/140  
78/78 [=====] - 39s 504ms/step - loss: 0.0709 - categorical\_acc  
uracy: 0.9786 - acc: 0.9786 - mean\_squared\_error: 0.0033  
Epoch 68/140  
78/78 [=====] - 38s 489ms/step - loss: 0.0673 - categorical\_acc  
uracy: 0.9780 - acc: 0.9780 - mean\_squared\_error: 0.0033  
Epoch 69/140  
78/78 [=====] - 38s 489ms/step - loss: 0.0629 - categorical\_acc  
uracy: 0.9816 - acc: 0.9816 - mean\_squared\_error: 0.0030  
Epoch 70/140  
78/78 [=====] - 38s 487ms/step - loss: 0.0719 - categorical\_acc  
uracy: 0.9758 - acc: 0.9758 - mean\_squared\_error: 0.0036  
Epoch 71/140  
78/78 [=====] - 39s 492ms/step - loss: 0.0665 - categorical\_acc  
uracy: 0.9796 - acc: 0.9796 - mean\_squared\_error: 0.0032  
Epoch 72/140  
78/78 [=====] - 38s 487ms/step - loss: 0.0619 - categorical\_acc  
uracy: 0.9790 - acc: 0.9790 - mean\_squared\_error: 0.0031  
Epoch 73/140  
78/78 [=====] - 41s 520ms/step - loss: 0.0649 - categorical\_acc  
uracy: 0.9800 - acc: 0.9800 - mean\_squared\_error: 0.0031  
Epoch 74/140  
78/78 [=====] - 40s 513ms/step - loss: 0.0673 - categorical\_acc  
uracy: 0.9796 - acc: 0.9796 - mean\_squared\_error: 0.0033  
Epoch 75/140  
78/78 [=====] - 39s 502ms/step - loss: 0.0668 - categorical\_acc  
uracy: 0.9802 - acc: 0.9802 - mean\_squared\_error: 0.0032  
Epoch 76/140  
78/78 [=====] - 38s 484ms/step - loss: 0.0621 - categorical\_acc  
uracy: 0.9828 - acc: 0.9828 - mean\_squared\_error: 0.0029  
Epoch 77/140  
78/78 [=====] - 38s 484ms/step - loss: 0.0673 - categorical\_acc  
uracy: 0.9792 - acc: 0.9792 - mean\_squared\_error: 0.0033  
Epoch 78/140  
78/78 [=====] - 37s 470ms/step - loss: 0.0634 - categorical\_acc  
uracy: 0.9810 - acc: 0.9810 - mean\_squared\_error: 0.0030  
Epoch 79/140  
78/78 [=====] - 39s 497ms/step - loss: 0.0686 - categorical\_acc  
uracy: 0.9794 - acc: 0.9794 - mean\_squared\_error: 0.0033  
Epoch 80/140  
78/78 [=====] - 39s 492ms/step - loss: 0.0659 - categorical\_acc  
uracy: 0.9810 - acc: 0.9810 - mean\_squared\_error: 0.0031  
Epoch 81/140  
78/78 [=====] - 37s 474ms/step - loss: 0.0578 - categorical\_acc  
uracy: 0.9824 - acc: 0.9824 - mean\_squared\_error: 0.0029  
Epoch 82/140  
78/78 [=====] - 38s 480ms/step - loss: 0.0614 - categorical\_acc  
uracy: 0.9810 - acc: 0.9810 - mean\_squared\_error: 0.0030  
Epoch 83/140  
78/78 [=====] - 38s 483ms/step - loss: 0.0605 - categorical\_acc  
uracy: 0.9824 - acc: 0.9824 - mean\_squared\_error: 0.0029  
Epoch 84/140  
78/78 [=====] - 39s 496ms/step - loss: 0.0668 - categorical\_acc  
uracy: 0.9796 - acc: 0.9796 - mean\_squared\_error: 0.0033

Epoch 85/140  
78/78 [=====] - 38s 484ms/step - loss: 0.0521 - categorical\_acc  
uracy: 0.9839 - acc: 0.9839 - mean\_squared\_error: 0.0026  
Epoch 86/140  
78/78 [=====] - 39s 496ms/step - loss: 0.0608 - categorical\_acc  
uracy: 0.9812 - acc: 0.9812 - mean\_squared\_error: 0.0029  
Epoch 87/140  
78/78 [=====] - 39s 491ms/step - loss: 0.0574 - categorical\_acc  
uracy: 0.9826 - acc: 0.9826 - mean\_squared\_error: 0.0029  
Epoch 88/140  
78/78 [=====] - 39s 493ms/step - loss: 0.0592 - categorical\_acc  
uracy: 0.9818 - acc: 0.9818 - mean\_squared\_error: 0.0028  
Epoch 89/140  
78/78 [=====] - 38s 486ms/step - loss: 0.0568 - categorical\_acc  
uracy: 0.9837 - acc: 0.9837 - mean\_squared\_error: 0.0027  
Epoch 90/140  
78/78 [=====] - 38s 482ms/step - loss: 0.0510 - categorical\_acc  
uracy: 0.9857 - acc: 0.9857 - mean\_squared\_error: 0.0024  
Epoch 91/140  
78/78 [=====] - 37s 476ms/step - loss: 0.0559 - categorical\_acc  
uracy: 0.9818 - acc: 0.9818 - mean\_squared\_error: 0.0028  
Epoch 92/140  
78/78 [=====] - 37s 471ms/step - loss: 0.0618 - categorical\_acc  
uracy: 0.9814 - acc: 0.9814 - mean\_squared\_error: 0.0030  
Epoch 93/140  
78/78 [=====] - 37s 468ms/step - loss: 0.0548 - categorical\_acc  
uracy: 0.9826 - acc: 0.9826 - mean\_squared\_error: 0.0027  
Epoch 94/140  
78/78 [=====] - 36s 465ms/step - loss: 0.0555 - categorical\_acc  
uracy: 0.9828 - acc: 0.9828 - mean\_squared\_error: 0.0026  
Epoch 95/140  
78/78 [=====] - 37s 470ms/step - loss: 0.0500 - categorical\_acc  
uracy: 0.9855 - acc: 0.9855 - mean\_squared\_error: 0.0025  
Epoch 96/140  
78/78 [=====] - 37s 467ms/step - loss: 0.0479 - categorical\_acc  
uracy: 0.9849 - acc: 0.9849 - mean\_squared\_error: 0.0023  
Epoch 97/140  
78/78 [=====] - 37s 466ms/step - loss: 0.0453 - categorical\_acc  
uracy: 0.9879 - acc: 0.9879 - mean\_squared\_error: 0.0020  
Epoch 98/140  
78/78 [=====] - 37s 469ms/step - loss: 0.0482 - categorical\_acc  
uracy: 0.9859 - acc: 0.9859 - mean\_squared\_error: 0.0024  
Epoch 99/140  
78/78 [=====] - 37s 466ms/step - loss: 0.0480 - categorical\_acc  
uracy: 0.9843 - acc: 0.9843 - mean\_squared\_error: 0.0024  
Epoch 100/140  
78/78 [=====] - 37s 469ms/step - loss: 0.0495 - categorical\_acc  
uracy: 0.9847 - acc: 0.9847 - mean\_squared\_error: 0.0024  
Epoch 101/140  
78/78 [=====] - 37s 466ms/step - loss: 0.0460 - categorical\_acc  
uracy: 0.9853 - acc: 0.9853 - mean\_squared\_error: 0.0023  
Epoch 102/140  
78/78 [=====] - 36s 467ms/step - loss: 0.0503 - categorical\_acc  
uracy: 0.9843 - acc: 0.9843 - mean\_squared\_error: 0.0025  
Epoch 103/140  
78/78 [=====] - 37s 467ms/step - loss: 0.0526 - categorical\_acc  
uracy: 0.9839 - acc: 0.9839 - mean\_squared\_error: 0.0026  
Epoch 104/140  
78/78 [=====] - 36s 463ms/step - loss: 0.0462 - categorical\_acc  
uracy: 0.9859 - acc: 0.9859 - mean\_squared\_error: 0.0022

```
Epoch 105/140
78/78 [=====] - 37s 465ms/step - loss: 0.0480 - categorical_acc
uracy: 0.9851 - acc: 0.9851 - mean_squared_error: 0.0023
Epoch 106/140
78/78 [=====] - 37s 472ms/step - loss: 0.0495 - categorical_acc
uracy: 0.9839 - acc: 0.9839 - mean_squared_error: 0.0024
Epoch 107/140
78/78 [=====] - 37s 467ms/step - loss: 0.0479 - categorical_acc
uracy: 0.9863 - acc: 0.9863 - mean_squared_error: 0.0024
Epoch 107: early stopping
```

```
In [17]: # Plot the progression of the accuracy through the epochs
import matplotlib.pyplot as plt
plt.plot(history.history['acc'])
plt.title('model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
In [179... # Demo to randomly pick a sign and demonstrate that it is predicted correctly

from random import randint
test_image = randint(0, len(data_train))
test = np.expand_dims(data_train[test_image], axis=0)
test = tf.keras.applications.inception_v3.preprocess_input(
    test, data_format=None
)
print(one_hot_translator(model.predict(test)))
plt.imshow(data_train[test_image])
```

```
1/1 [=====] - 0s 37ms/step
Stop
```

```
Out[179... <matplotlib.image.AxesImage at 0x248214161c0>
```






```
In [19]: # Calculate the predictions for the test values
         predictions = model.predict(x_test)
```

39/39 [=====] - 9s 231ms/step

```
In [20]: evaluation = model.evaluate(x_test, t_test)
         print("Test run accuracy is {}".format(evaluation[-1]))
```

39/39 [=====] - 9s 228ms/step - loss: 0.3067 - categorical\_accuracy: 0.9379 - acc: 0.9379 - mean\_squared\_error: 0.0100  
Test run accuracy is 0.010019068606197834

```
In [21]: # Create a Confusion Matrix to show the weakness in the model
         predicted_values = translate_all(predictions)
         real_values = translate_all(t_test)
         from sklearn.metrics import confusion_matrix
         from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
         cfm = confusion_matrix(real_values, predicted_values)
         disp = ConfusionMatrixDisplay(confusion_matrix=cfm, display_labels=['Stop', 'Yield', 'Red
         disp.plot()
         plt.show()
```



```
In [22]: # Test Run Data
```

```
# Testing optimizers
# NADAM=                                times = [275,272,286], Loss after 3 runs = [
# ADAM=                                times = [190,186,190], Loss after 3 runs = [
# SGD(.001, Nesterov=No, Momentum = No), times = [182,180,181], Loss after 3 runs = [
# SGD(.01, Nesterov=No, Momentum = No),  times = [180,180,179], Loss after 3 runs = [
# SGD(.01, Nesterov=Yes, Momentum = 0.25), times = [183,183,183], Loss after 3 runs = [
# SGD(.01, Nesterov=Yes, Momentum = 0.5), times = [189,183,183], Loss after 3 runs = [
# SGD(.01, Nesterov=Yes, Momentum = 0.75), times = [185,185,184], Loss after 3 runs = [
```


In [97]:

```
# This method is used to show an example of the post processed test data
from random import randint
import matplotlib.pyplot as plt
test_image = randint(0,len(x_train))
plt.imshow(x_train[test_image])
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Out[97]:

<matplotlib.image.AxesImage at 0x24821a2caf0>



In [16]:

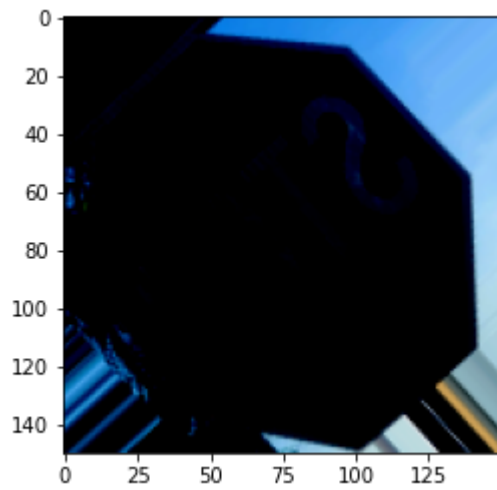
```
# Original 140 Epoch run was done with .25-.75 brightness range
```

In [37]:

```
# This method is used to show an example of the post processed data
from random import randint
x,y = next(train_generator)

plt.imshow(x[randint(0,63)])
plt.show()
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



In [253...

```
Load_image = tf.keras.preprocessing.image.load_img('New.jpg')
og_image = Load_image.copy()
Load_image = np.array(tf.cast(tf.image.resize(Load_image,(150,150)), np.uint8))
print(Load_image.shape)
Load_image = tf.keras.applications.inception_v3.preprocess_input(
    Load_image, data_format=None
)
Load_image = tf.expand_dims(Load_image,0)
prediction = model.predict(Load_image)
print(prediction)
print(one_hot_translator(prediction))
plt.imshow(og_image)
calculate_likelyhood(prediction)
plt.plot([0,1,2,3,4,5,6,7,8,9],prediction[0])
plt.show()
```


(150, 150, 3)

1/1 [=====] - 0s 38ms/step

```
[[2.3534343e-05 1.6287031e-05 6.8655545e-03 3.1475363e-06 9.8857528e-01
 6.3804933e-04 1.9053403e-05 1.5509835e-05 2.7066142e-07 3.8431643e-03]]
```

Roundabout





```
In [250... def calculate_std(np_array):  
    translated_values = []  
    for i in np_array:  
        translated_values.append((np.argmax(i))/sum(i))  
    return np.array(translated_values)
```

```
In [251... x = calculate_std(predictions)  
print(x)  
print(np.argmax(x))
```

```
[2.00000003 7.00000062 6.99999951 ... 2.00000004 5.99999922 6.99999949]  
422
```

```
In [ ]:
```